

```
# Bike Sharing System -Part I: Exploratory Analysis
```

```
# import the necessary packages
```

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn import linear_model
import csv
import os
import pandas as pd
import shutil, glob
```

```
# After downloading the train.csv and test.csv files, we are reading the two files in DataFrames
```

```
path='F:/dataproyects'
```

```
df = pd.read_csv('F:/dataproyects/train.csv')
dftest=pd.read_csv('F:/dataproyects/test.csv')
```

```
# Exploratory analysis of the train DataFrame
```

```
# In my case, there where some problems with identifying the 'count' column of the DataFrame, so we renamed it
```

```
#Rename the last column of the DataFrame
```

```
df.rename(columns={'count': 'totalcustomers'}, inplace=True)
list(df.columns.values)
```

```
['datetime',
 'season',
 'holiday',
 'workingday',
 'weather',
 'temp',
 'atemp',
 'humidity',
 'windspeed',
 'casual',
 'registered',
 'totalcustomers']
```

```

#feature engineering
# We want to analyze the bike sharing distribution on hours, days, months, years.
# so Let's regain these informations from the 'datetime' column
df['datetime']=pd.to_datetime(df['datetime'])
print(df['datetime'].dtype)
df['year'] = df['datetime'].dt.year
df['month'] = df['datetime'].dt.month
df['day'] = df['datetime'].dt.day
df['hour']=df['datetime'].dt.hour
cols = df.columns.tolist()
cols=cols[-4:]+cols[:-4]
df=df[cols]
df=df.drop('datetime', axis=1)
list(df.columns.values)

```

```
datetime64[ns]
```

```

['year',
 'month',
 'day',
 'hour',
 'season',
 'holiday',
 'workingday',
 'weather',
 'temp',
 'atemp',
 'humidity',
 'windspeed',
 'casual',
 'registered',
 'totalcustomers']

```

```

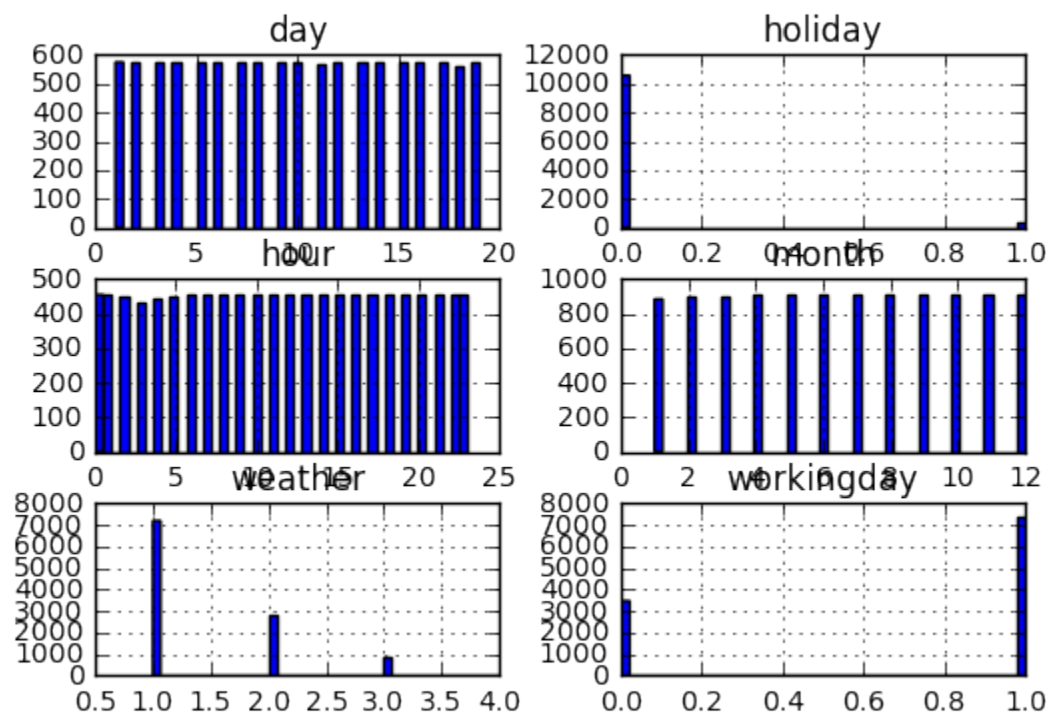
# check if there are any missing values in the DataFrame: in this case there are no missing values in our database
# the isnull() detects missing values in the specified DataFrame df
df.isnull().any()
# there are 10886 entries in our dataframe
len(df.weather)

```

```
10886
```

```
# Analyze the data in order to bring a better understanding of the data we are looking at.
# The shape and dispersion of the data output can help significantly in noticing the evolution of one variable.
# Visualisation tools are extremely valuable here
# For each of the predictor variables, we first see the histogram :
# (season, workingday, weather, temp, atemp, humidity, windspeed, casual, registered, count).
# We notice that most of our variables do not follow a normal distribution
df1=df[['month', 'day', 'hour', 'holiday','weather','workingday']]
df2=df[['temp', 'atemp','humidity', 'windspeed']]
df3=df[['casual', 'registered', 'totalcustomers']]
```

```
df1.hist(layout=(3,2), bins='rice')
plt.show()
# nb of bins has to be manually chosen, to reflect the nb of elements in the bin
```



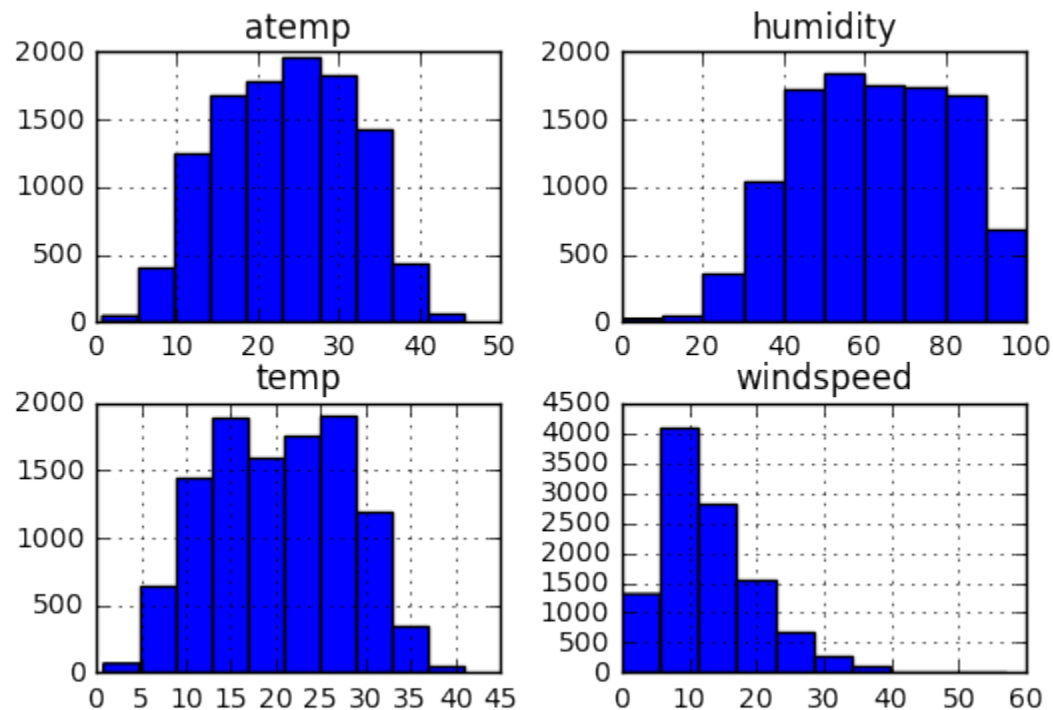
png

```
print len(df[df['day']==1])
print len(df[df['weather']==1])
```

```
575
7192
```

```
# Histograms above correspond to variables of nominal type data.
# Nominal refers to data that is categorical, for example subtracting one month from
another has no meaning.
# workingday predictor: if day is neither weekend nor holiday is 1, otherwise is 0.
# weather predictor: + weathersit :
#- 1: Clear, Few clouds, Partly cloudy, Partly cloudy
#- 2: Mist + Cloudy, Mist + Broken clouds, Mist + Few clouds, Mist
#- 3: Light Snow, Light Rain + Thunderstorm + Scattered clouds, Light Rain + Scattered clouds
#- 4: Heavy Rain + Ice Pallets + Thunderstorm + Mist, Snow + Fog
# temperature predictor: Normalized temperature in Celsius.
# The values are derived via (t-t_min)/(t_max-t_min), t_min=-8, t_max=+39 (only in hourly scale)
# hour predictor
# The most values seems to be collected in the first and last hour of the day
```

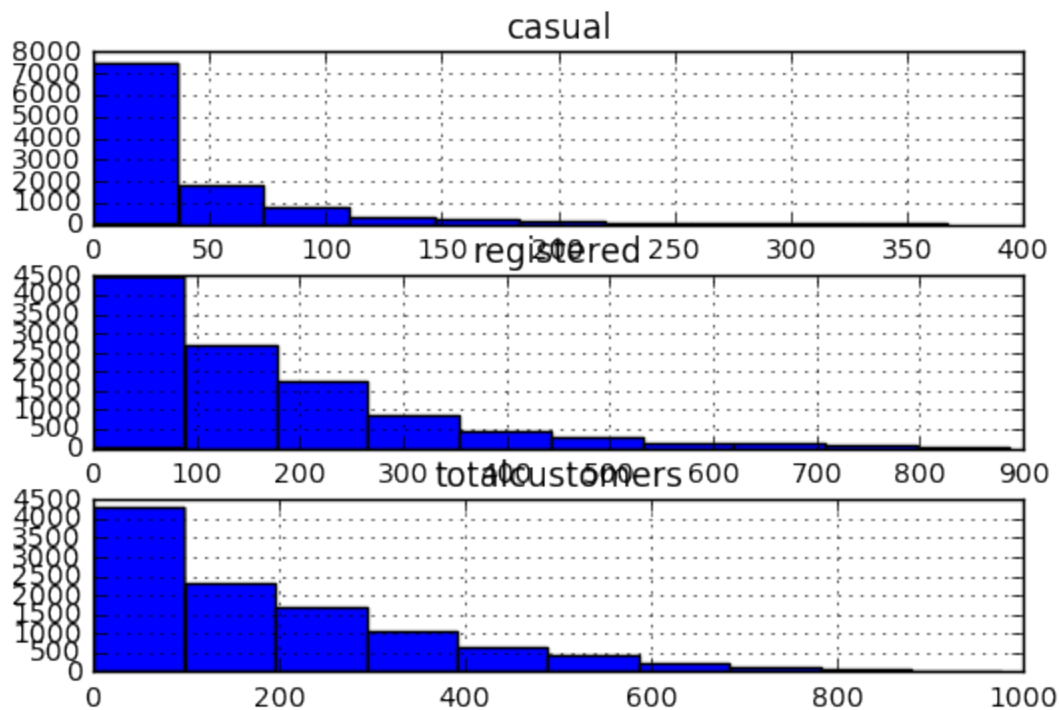
```
df2.hist(layout=(2,2))
plt.show()
```



png

```
# The histograms above are example of interval data, these type of data have scale.
# The variables do not follow a normal distribution either.
```

```
df3.hist(layout=(3,1))
plt.show()
```



png

The diagrams are positively skewed

Descriptive statistic elements

Let's evaluate the mean and the median values of the customers for each season.

We check for noticeable differences between the mean and the median.

As the mean is more susceptible to outliers, it would conceivably be distorted greatly in the presence of a large number of outliers or large outliers values.

For variables depicting the total number of clients, casual or registered show differences between the

mean value and the median. This shows the possible existence of outliers.

```
df["casual"].mean()
```


```
36.02195480433584
```

```
df["casual"].median()
```

```
17.0
```

```
# casual customers boxplot
plt.boxplot(df.casual)
plt.show()
data=df.casual
median = np.median(data)
upper_quartile = np.percentile(data, 75)
lower_quartile = np.percentile(data, 25)
iqr = upper_quartile - lower_quartile
upper_whisker = data[data<=upper_quartile+1.5*iqr].max()
lower_whisker = data[data>=lower_quartile-1.5*iqr].min()

# the Boxplot does show the presence of many outliers
```



png

```
df["totalcustomers"].mean()
```

```
191.57413191254824
```


```
df["totalcustomers"].median()
```

```
145.0
```

```
# for the variables like season, holiday, workingday, a median or mean value makes no sense.
```

```
# boxes: the main body of the boxplot showing the quartiles and the median and confidence intervals if enabled.
# medians: horizontal lines at the median of each box.
# whiskers: the vertical lines extending to the most extreme, non-outlier data points.
# caps: the horizontal lines at the ends of the whiskers.
# fliers: points representing data that extend beyond the whiskers (outliers).
```

```
df3.boxplot()
plt.show()
```



png

```
# we can then extract all the information in the boxplots that you are interested in,
# e.g. median, upper_quartile, iqr, etc.
# I have wrote an example for the 'totalcustomers' variable
data=df.totalcustomers
median = np.median(data)
upper_quartile = np.percentile(data, 75)
lower_quartile = np.percentile(data, 25)
iqr = upper_quartile - lower_quartile
upper_whisker = data[data<=upper_quartile+1.5*iqr].max()
lower_whisker = data[data>=lower_quartile-1.5*iqr].min()
print upper_whisker
print iqr
```

```
647
242.0
```

For the humidity variable there seems to be no significant difference between the mean and the median values.

Same situation for the windspeed and for the temperature, which means that most probably outliers are not present

```
df["humidity"].mean()
```

61.88645967297446

```
df["humidity"].median()
```

62.0

```
df["windspeed"].mean()
```

12.799395406945093

```
df["windspeed"].median()
```

12.998

```
df["temp"].mean()
```

20.230859819952173

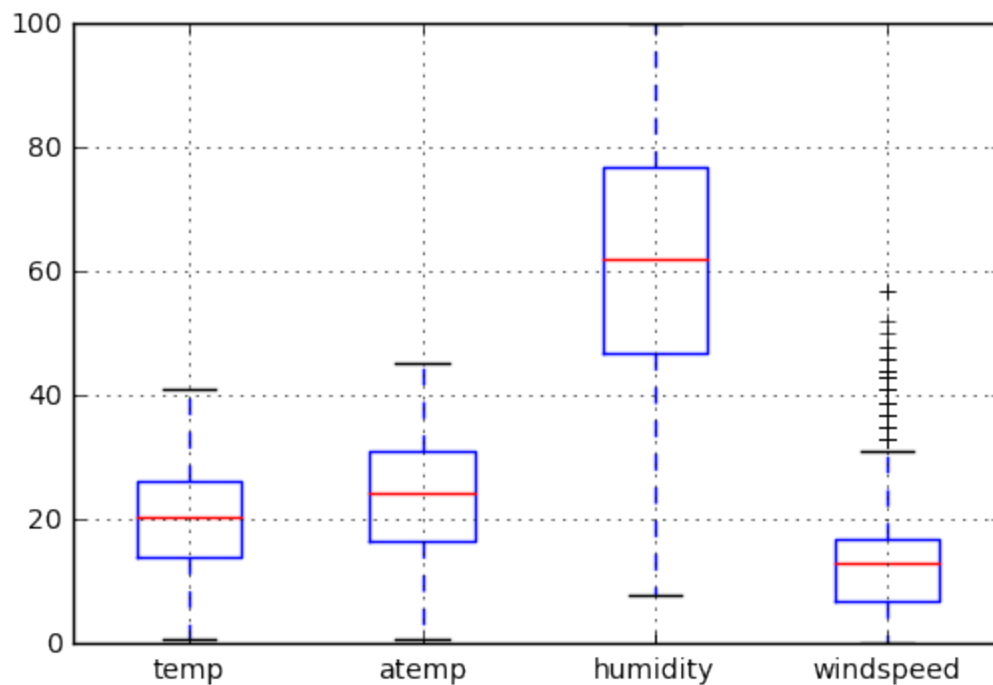
```
df["temp"].median()
```

20.5

```
df2.boxplot()
```

```
plt.show()
```

for the variable windspeed a small number of outliers is present, the influence on the mean value is not significantly large



png

```
# We will continue to investigate these attributes in the nex paragraph about normalit  
y  
# The Shapiro-Wills normality test
```

```
# Let us start by checking the correlation between our variables  
# Analyze the level of correlation between the DataFrame variables
```

```
df.corr(method='pearson', min_periods=1)  
# The correlation matrix show some existing relation between the nb of casual and tota  
l customers,  
# between the registered and total nb of customers,  
# however, weather characteristic does not seems to have a significant impact on the n  
b of customers
```

```
<tr style="text-align: right;">
  <th></th>
  <th>year</th>
  <th>month</th>
  <th>day</th>
  <th>hour</th>
  <th>season</th>
  <th>holiday</th>
  <th>workingday</th>
  <th>weather</th>
  <th>temp</th>
  <th>atemp</th>
  <th>humidity</th>
  <th>windspeed</th>
  <th>casual</th>
  <th>registered</th>
  <th>totalcustomers</th>
</tr>
```

```
<tr>
  <th>year</th>
  <td>1.000000</td>
  <td>-0.004932</td>
  <td>0.001800</td>
  <td>-0.004234</td>
  <td>-0.004797</td>
  <td>0.012021</td>
  <td>-0.002482</td>
  <td>-0.012548</td>
  <td>0.061226</td>
  <td>0.058540</td>
  <td>-0.078606</td>
  <td>-0.015221</td>
  <td>0.145241</td>
  <td>0.264265</td>
  <td>0.260403</td>
```

```
</tr>
```

```
<tr>
  <th>month</th>
  <td>-0.004932</td>
  <td>1.000000</td>
  <td>0.001974</td>
  <td>-0.006818</td>
  <td>0.971524</td>
  <td>0.001731</td>
  <td>-0.003394</td>
  <td>0.012144</td>
  <td>0.257589</td>
  <td>0.264173</td>
  <td>0.204537</td>
  <td>-0.150192</td>
  <td>0.092722</td>
  <td>0.169451</td>
  <td>0.166862</td>
```

```
</tr>
```

```
<tr>
  <th>day</th>
  <td>0.001800</td>
  <td>0.001974</td>
  <td>1.000000</td>
  <td>0.001132</td>
  <td>0.001729</td>
  <td>-0.015877</td>
  <td>0.009829</td>
  <td>-0.007890</td>
  <td>0.015551</td>
  <td>0.011866</td>
```

```
<td>-0.011335</td>
<td>0.036157</td>
<td>0.014109</td>
<td>0.019111</td>
<td>0.019826</td>
</tr>
<tr>
<th>hour</th>
<td>-0.004234</td>
<td>-0.006818</td>
<td>0.001132</td>
<td>1.000000</td>
<td>-0.006546</td>
<td>-0.000354</td>
<td>0.002780</td>
<td>-0.022740</td>
<td>0.145430</td>
<td>0.140343</td>
<td>-0.278011</td>
<td>0.146631</td>
<td>0.302045</td>
<td>0.380540</td>
<td>0.400601</td>
</tr>
<tr>
<th>season</th>
<td>-0.004797</td>
<td>0.971524</td>
<td>0.001729</td>
<td>-0.006546</td>
<td>1.000000</td>
<td>0.029368</td>
<td>-0.008126</td>
<td>0.008879</td>
<td>0.258689</td>
<td>0.264744</td>
<td>0.190610</td>
<td>-0.147121</td>
<td>0.096758</td>
<td>0.164011</td>
<td>0.163439</td>
</tr>
<tr>
<th>holiday</th>
<td>0.012021</td>
<td>0.001731</td>
<td>-0.015877</td>
<td>-0.000354</td>
<td>0.029368</td>
```

```
<td>1.000000</td>
<td>-0.250491</td>
<td>-0.007074</td>
<td>0.000295</td>
<td>-0.005215</td>
<td>0.001929</td>
<td>0.008409</td>
<td>0.043799</td>
<td>-0.020956</td>
<td>-0.005393</td>
</tr>
<tr>
<th>workingday</th>
<td>-0.002482</td>
<td>-0.003394</td>
<td>0.009829</td>
<td>0.002780</td>
<td>-0.008126</td>
<td>-0.250491</td>
<td>1.000000</td>
<td>0.033772</td>
<td>0.029966</td>
<td>0.024660</td>
<td>-0.010880</td>
<td>0.013373</td>
<td>-0.319111</td>
<td>0.119460</td>
<td>0.011594</td>
</tr>
<tr>
<th>weather</th>
<td>-0.012548</td>
<td>0.012144</td>
<td>-0.007890</td>
<td>-0.022740</td>
<td>0.008879</td>
<td>-0.007074</td>
<td>0.033772</td>
<td>1.000000</td>
<td>-0.055035</td>
<td>-0.055376</td>
<td>0.406244</td>
<td>0.007261</td>
<td>-0.135918</td>
<td>-0.109340</td>
<td>-0.128655</td>
</tr>
<tr>
<th>temp</th>
```

```
<td>0.061226</td>
<td>0.257589</td>
<td>0.015551</td>
<td>0.145430</td>
<td>0.258689</td>
<td>0.000295</td>
<td>0.029966</td>
<td>-0.055035</td>
<td>1.000000</td>
<td>0.984948</td>
<td>-0.064949</td>
<td>-0.017852</td>
<td>0.467097</td>
<td>0.318571</td>
<td>0.394454</td>
</tr>
<tr>
<th>atemp</th>
<td>0.058540</td>
<td>0.264173</td>
<td>0.011866</td>
<td>0.140343</td>
<td>0.264744</td>
<td>-0.005215</td>
<td>0.024660</td>
<td>-0.055376</td>
<td>0.984948</td>
<td>1.000000</td>
<td>-0.043536</td>
<td>-0.057473</td>
<td>0.462067</td>
<td>0.314635</td>
<td>0.389784</td>
</tr>
<tr>
<th>humidity</th>
<td>-0.078606</td>
<td>0.204537</td>
<td>-0.011335</td>
<td>-0.278011</td>
<td>0.190610</td>
<td>0.001929</td>
<td>-0.010880</td>
<td>0.406244</td>
<td>-0.064949</td>
<td>-0.043536</td>
<td>1.000000</td>
<td>-0.318607</td>
<td>-0.348187</td>
```

```
<td>-0.265458</td>
<td>-0.317371</td>
</tr>
<tr>
<th>windspeed</th>
<td>-0.015221</td>
<td>-0.150192</td>
<td>0.036157</td>
<td>0.146631</td>
<td>-0.147121</td>
<td>0.008409</td>
<td>0.013373</td>
<td>0.007261</td>
<td>-0.017852</td>
<td>-0.057473</td>
<td>-0.318607</td>
<td>1.000000</td>
<td>0.092276</td>
<td>0.091052</td>
<td>0.101369</td>
</tr>
<tr>
<th>casual</th>
<td>0.145241</td>
<td>0.092722</td>
<td>0.014109</td>
<td>0.302045</td>
<td>0.096758</td>
<td>0.043799</td>
<td>-0.319111</td>
<td>-0.135918</td>
<td>0.467097</td>
<td>0.462067</td>
<td>-0.348187</td>
<td>0.092276</td>
<td>1.000000</td>
<td>0.497250</td>
<td>0.690414</td>
</tr>
<tr>
<th>registered</th>
<td>0.264265</td>
<td>0.169451</td>
<td>0.019111</td>
<td>0.380540</td>
<td>0.164011</td>
<td>-0.020956</td>
<td>0.119460</td>
<td>-0.109340</td>
```

```

<td>0.318571</td>
<td>0.314635</td>
<td>-0.265458</td>
<td>0.091052</td>
<td>0.497250</td>
<td>1.000000</td>
<td>0.970948</td>
</tr>
<tr>
<th>totalcustomers</th>
<td>0.260403</td>
<td>0.166862</td>
<td>0.019826</td>
<td>0.400601</td>
<td>0.163439</td>
<td>-0.005393</td>
<td>0.011594</td>
<td>-0.128655</td>
<td>0.394454</td>
<td>0.389784</td>
<td>-0.317371</td>
<td>0.101369</td>
<td>0.690414</td>
<td>0.970948</td>
<td>1.000000</td>
</tr>

```

The correlation coefficient between the temp and atemp variables is close to 1, i.e. 0.984948. The same situation for registered and total customer number.

In order to avoid a multicollinearity situation, we will eliminate the atemp variable from our dataframe

season and month very high correlation 0.97


the correlation is moderate also between the casual and total customers 0.69

This correlation can also be seen in the figures below

```


plt.scatter(df.temp, df.atemp)
plt.show()

```


png

```
plt.scatter(df.registered, df.totalcustomers)
plt.show()
```




png

```
plt.scatter(df.casual, df.totalcustomers)
plt.show()
# moderate correlation
```



png

```
plt.scatter(df.season, df.month)
plt.show()
```



png

```
# Heatmap to see the correlation matrix on a different Level

# import seaborn as sns
# sns.heatmap(df[['year', 'month', 'day', 'hour', 'holiday', 'workingday', 'weather', 'temp', 'humidity', 'windspeed', 'casual', 'totalcustomers']].corr(), annot=True)
# plt.show()
# There is an issue with matplotlib boxplot fliers not showing up when seaborn is imported,
# even when fliers are explicitly enabled. In these conditions this heatmap, showing in a nice way the correlation coefficients is here commented.
# you can use it, but pay attention that in the boxplot histogram, you will not be able to see the fliers.
```

```
from scipy import stats
```

```
# Perform the Shapiro-Wilk test for normality.
# Some of our variables have a distribution that could be normal.
# The Shapiro-Wilk test is in fact a test for the existence of a normal distribution in data: it tests the null hypothesis that data are normal.
# If the p-value is greater than the chosen alpha level, then the null hypothesis that the data came from a normally distributed population cannot be rejected (e.g., for an alpha level of 0.05, a data set with a p-value of 0.02 rejects the null hypothesis that the data are from a normally distributed population).[2] However, since the test is biased by sample size,[3] the test may be statistically significant from a normal distribution in any large samples. Thus a Q-Q plot is required for verification in addition to the test.
# The Shapiro-Wilk test tests the null hypothesis that the data was drawn from a normal distribution: if the p-value is less than the chosen alpha level (0.05 here),
# then the null hypothesis is rejected and there is evidence that the data tested are not from a normally distributed population.
# The test shows that the data are actually not normally distributed. The p-values are extremely small.
newdf=df.drop('atemp',axis=1).drop('season',axis=1)
stats.shapiro(newdf.humidity)
```

```
C:\Users\ss_cr\Anaconda1\lib\site-packages\scipy\stats\morestats.py:1326: UserWarning:
p-value may not be accurate for N > 5000.
```

```
warnings.warn("p-value may not be accurate for N > 5000.")
```

```
(0.9822689294815063, 1.245496990918048e-34)
```

```
# because of this message, we also back our normality test with a qqplot
```

```
stats.shapiro(newdf.temp)
```

```
(0.9804092645645142, 4.47221826500091e-36)
```

```
stats.shapiro(newdf.windspeed)
```

```
(0.9587375521659851, 0.0)
```


```
stats.shapiro(newdf.totalcustomers)
```

```
(0.8783667087554932, 0.0)
```

```
stats.shapiro(newdf.casual)
```


```
(0.7056357264518738, 0.0)
```

```
# the Shapiro Wills test is biased by sample size, our data based has more than 5000 e  
ntries  
# the test may be statistically significant from a normal distribution in any large sa  
mles.  
# Thus a QâQ plot is required for verification in addition to the test.  
# test the normal distribution case with quantile - quantile plot with scipy  
# Th thick blue line represents the distribution of the actual variable from the data  
set  
# and the straight red line is a mapping of what the normal distribution would look li  
ke  
import pylab  
stats.probplot(newdf.casual, dist="norm", plot=pylab)  
pylab.show()  
# clearly not normally distributed
```




png

```
stats.probplot(newdf.registered, dist="norm", plot=pylab)
pylab.show()
```




png

```
stats.probplot(newdf.hour, dist="norm", plot=pylab)
pylab.show()
```




png

```
stats.probplot(newdf.month, dist="norm", plot=pylab)
pylab.show()
```



png

```
stats.probplot(newdf.humidity, dist="norm", plot=pylab)
pylab.show()
```

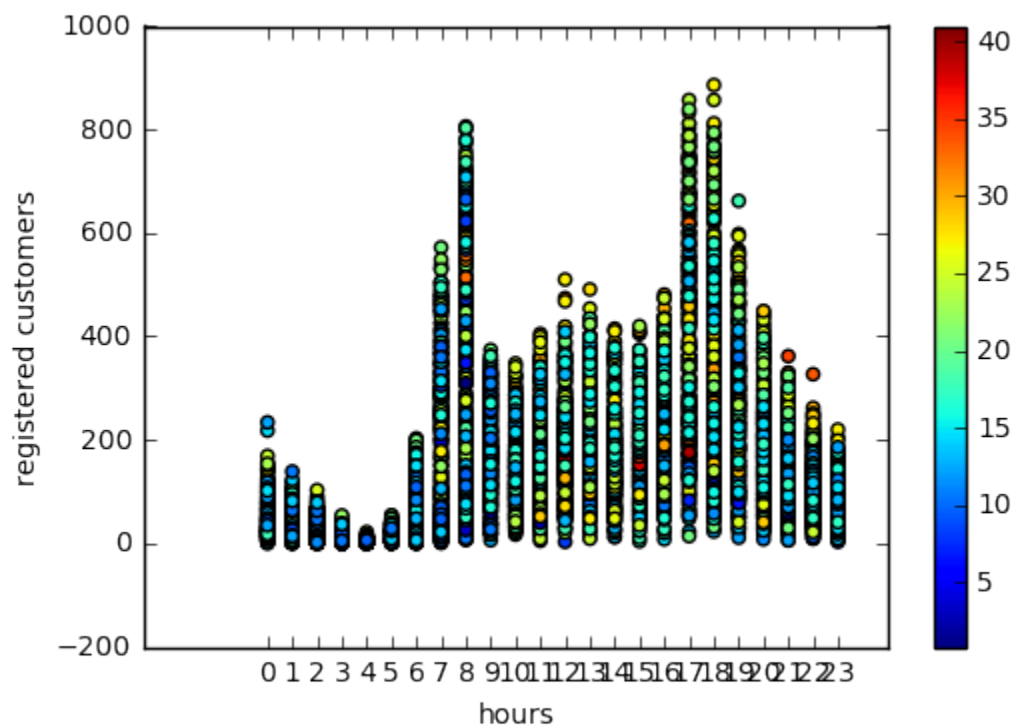


png

```
# It looks like the tested variables are not normally distributed
```

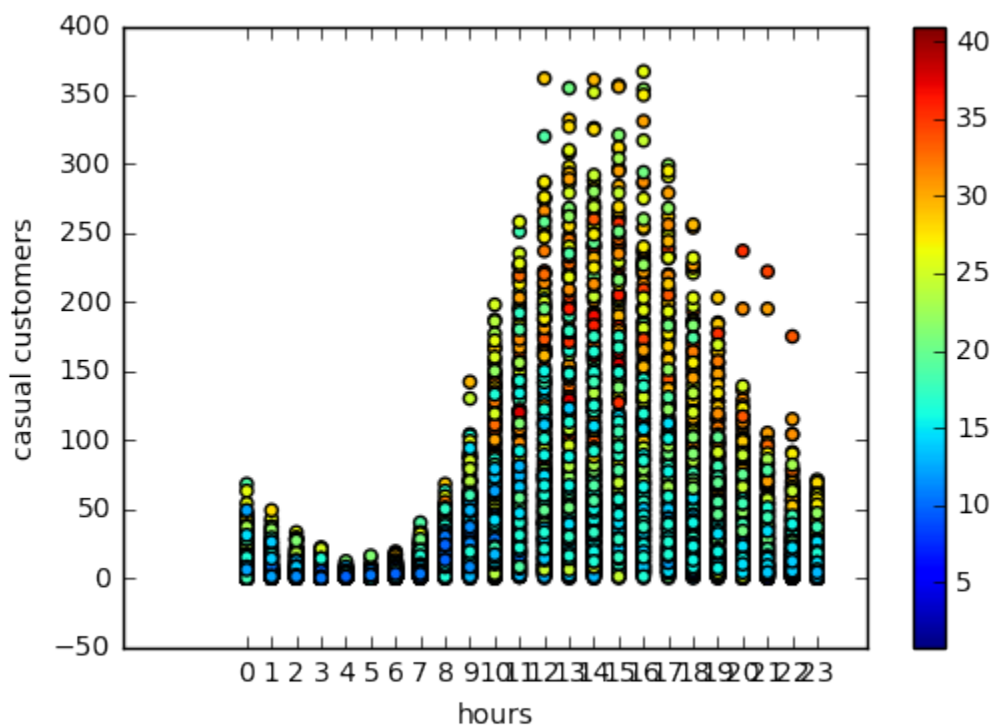
```
# Let's use some visualisation tools to explore our data
```

```
plt.scatter(newdf.hour, newdf.registered, c=df.temp,cmap='jet')
plt.colorbar()
ticks=np.arange(0,24,1)
labels = range(ticks.size)
plt.xticks(ticks,labels)
plt.xlabel('hours')
plt.ylabel('registered customers')
plt.show()
# It looks like the largest amount of registered customers is between 7 and 8, when pe
# ople probably go to work
# and then again between 17.00 and 18.00, when people leave work.
# This diagram clearly makes a lot of sense.
```



png



```
plt.scatter(newdf.hour, newdf.casual, c=df.temp, cmap='jet')
plt.colorbar()
ticks=np.arange(0,24,1)
labels = range(ticks.size)
plt.xticks(ticks,labels)
plt.xlabel('hours')
plt.ylabel('casual customers')
plt.show()
# It Looks Like the nb of casual customers (most of them probably tourists) starts to i
ncrease around 9.00 a.m. and then decreases agin starting with 22.00.
# The number of casul customers is clearly larger for optimal wheater between 20 and 3
0 degrees.
```



png

```
labels_hour=df.hour.unique()
values_hour=df.groupby('hour')['casual'].sum()


plt.pie(values_hour, labels=labels_hour,
        autopct='%1.1f%%', shadow=False, startangle=90)
plt.show()
# Same conclusion as above: the number of customers increase after 11a.m. and decreas
e after 6p.m.
```



png

```
labels_hour=df.hour.unique()
values_hour=df.groupby('hour')['totalcustomers'].sum()

plt.pie(values_hour, labels=labels_hour,
        autopct='%1.1f%%', shadow=False, startangle=90)
plt.show()
# We regain the conclusion from figures above: totalcustomers/registered number is lar
ger around 8 a.m and 5-6 p.m.
```




png

```

labels_month = ['Jan','Feb','March','April', 'May', 'June', 'July', 'Aug', 'Sept', 'Oct', 'Nov', 'Dec']
values_month = df.groupby('month')['totalcustomers'].sum()
plt.pie(values_month, labels=labels_month,
        autopct='%1.1f%%', shadow=False, startangle=90)
plt.show()
# The number of totalcustomers is larger between May and October

```




png

```

labels_year=df.year.unique()
values_year=df.groupby('year')['totalcustomers'].sum()


plt.pie(values_year, labels=labels_year,
        autopct='%1.1f%%', shadow=False, startangle=90)
plt.show()
# This shows that the number of bike demands is much larger for 2012

```




png

```
labels_day=df.workingday.unique()
values_day=df.groupby('workingday')['casual'].sum()
plt.pie(values_day, labels=labels_day,
        autopct='%1.1f%%', shadow=False, startangle=90)
plt.show()
# slight increase in the number of casual bikers over the weekends
```




png

```
plt.scatter(newdf.hour, newdf.casual, c=df.weather)
plt.colorbar()
ticks=np.arange(0,12,1)
labels = range(ticks.size)
plt.xticks(ticks,labels)
plt.xlabel('hours')
plt.ylabel('casual customers')
plt.show()
# casual customers number has the same hourly distribution in terms of wheater as in t
erms of temperature
```




png

```
plt.scatter(newdf.hour, newdf.totalcustomers, c=df.windspeed, cmap='jet')
plt.colorbar()
ticks=np.arange(0,24,1)
labels = range(ticks.size)
plt.xticks(ticks,labels)
plt.xlabel('hours')
plt.ylabel('registered customers')
plt.show()
# the larger numb of customers use the bike for wind temperature up to 18km/h
```




png

```
plt.scatter(newdf.humidity, newdf.totalcustomers)
plt.show()
```



png

```
plt.scatter(newdf.windspeed, newdf.totalcustomers)
plt.show()
# the wind speed has a certain influence on the nb of customers
```



png

```
#Bike sharing system - Part II - Building the model ...
```

```
# Building the regression model. First we split the dataframe into a training and test
case (60%, 40%).
# test_size=0.4 inside the function indicates the percentage of the data that should b
e held over for testing.
```

```
from sklearn.cross_validation import train_test_split
training, testing= train_test_split( newdf, test_size=0.4, random_state=1 )
print len(training)
print len(testing)
```

6531

4355

```
C:\Users\ss_cr\Anaconda1\lib\site-packages\sklearn\cross_validation.py:44: Deprecation
Warning: This module was deprecated in version 0.18 in favor of the model_selection mo
dule into which all the refactored classes and functions are moved. Also note that th
e interface of the new CV iterators are different from that of this module. This modul
e will be removed in 0.20.
```

```
"This module will be removed in 0.20.", DeprecationWarning)
```

```
# ...
```